

# Supplementary Information

## Supplementary Methods

### Representative transcript selection from Ribo-seq alignments

For each cellular context, exon annotations were parsed from the GENCODE GTF file to define non-overlapping exon bins from the union of transcript exon boundaries, gene intervals spanning each annotated gene and splice junctions connecting consecutive exons. MANE Select tags were extracted to identify representative isoforms. Ribo-seq reads were counted genome-wide at their 5' ends, and junction-spanning reads were identified from N operations in the CIGAR string. Exon bins supported by at least two Ribo-seq reads were designated informative features; observed splice junctions provided additional isoform support.

For each gene, a binary feature-by-transcript matrix linked informative features to compatible transcript isoforms. Redundant non-MANE transcripts with identical feature sets and nested non-MANE transcripts whose feature sets were strict subsets of another transcript were removed. An iterative greedy procedure then selected the transcript explaining the largest number of informative features not yet covered. When multiple transcripts shared the same informative features, preference was given to the transcript containing fewer interspersed low-support features. Selection continued until no additional transcript increased informative-feature coverage. MANE Select transcripts were retained irrespective of whether they added a new informative feature.

### Nuclease-aware P-site inference

To infer P-sites from Ribo-seq libraries generated under different nuclease conditions, including RNase I and micrococcal nuclease (MNase), we developed a nuclease-aware LightGBM classifier (version 4.6.0) that assigned a read-specific offset from the RPF 5' end. For each nuclease class, informative fragment lengths and initial offset ranges were identified from read-density peaks around annotated translation initiation or termination sites and from CDS three-nucleotide periodicity. Candidate labels were the offsets between an RPF 5' end and the corresponding annotated reference site.

For each read, the classifier used one-hot-encoded nucleotide windows spanning 8 nucleotides on either side of both the 5' and 3' ends, together with a one-hot representation of fragment length, to predict the candidate offset as a multiclass outcome. Models were fitted

with a learning rate of 0.1, up to 2,000 trees, class-balanced weights and early stopping after 100 rounds without improvement. Fragment lengths with strong aggregate periodicity (periodicity score at least 0.55) were assigned a fixed, frame-adjusted offset. For all other informative lengths, the P-site was assigned using the highest-probability offset within the length-specific candidate range. Read counts assigned to the same transcript position were summed to obtain the nucleotide-resolution P-site count profile.

### **Selection of informative transcript–context profiles**

After P-site inference, transcript–context profiles were required to have matched RNA-seq abundance of at least 1.0 reads per million mapped reads (RPM) and sufficient P-site signal. For transcripts with valid annotated CDS coordinates and a CDS longer than 45 nucleotides, P-site depth and coverage were calculated over the CDS. CDS depth was defined as the total number of inferred P-sites divided by CDS length, and CDS coverage as the fraction of consecutive CDS codons containing at least one inferred P-site. CDS P-site depth and coverage were each required to be at least 0.3, and more than 33% of CDS P-sites were required to fall in frame 0.

For transcripts without a valid annotated CDS, P-site depth and coverage were calculated over the full transcript. Full-length depth was defined as the total number of inferred P-sites divided by transcript length, and full-length coverage as the fraction of consecutive three-nucleotide bins containing at least one inferred P-site. Full-length P-site depth and coverage were each required to be at least 0.1; no frame-periodicity criterion was applied.

### **RNA-abundance normalization and target construction**

We transformed each P-site count profile into an RNA-abundance-adjusted relative ribosome occupancy profile. Within each transcript–context profile, non-zero P-site counts were winsorized at the 99.5th percentile. To estimate transcript-level translation relative to RNA abundance, total RPF and RNA-seq counts were each augmented by a pseudocount of 10 and centered-log-ratio (CLR) transformed across retained transcripts within the same cellular context. An ordinary least-squares regression of CLR-transformed RPF counts on CLR-transformed RNA-seq counts was then fitted. The residual,  $r_t$ , represented translation not explained by measured RNA abundance, and  $\exp(r_t)$  was used as a multiplicative translation-efficiency scale. For transcript  $t$  and nucleotide position  $\ell$ , the final target was defined as:

$$y_{t\ell} = \log \left[ 1 + \frac{c_{t\ell}^{(w)}}{\hat{c}_t^{(+)}} \exp(r_t) \right]$$

$$\hat{c}_t^{(+)} = \frac{1}{|P_t^+|} \sum_{l \in P_t^+} c_{tl}^{(w)}, r_t = \text{clr}(R_t + 10) - [\beta_0 + \beta_1 \text{clr}(A_t + 10)].$$

Here  $c_{ti}^{(w)}$  denotes the winsorized P-site count,  $P_t^+$  is the set of positions with a positive winsorized count,  $\bar{c}_t^{(+)}$  is the mean over those positions,  $R_t$  is the transcript-level RPF count and  $A_t$  is the matched RNA-seq count. The regression coefficients  $\beta_0$  and  $\beta_1$  were estimated separately within each cellular context. The resulting  $y_{ti}$  is dimensionless and represents a relative occupancy signal after adjustment for transcript abundance; it is not an absolute ribosome density, a single-molecule measurement or a site-occupancy probability. This procedure yielded 1.8 million transcript–context profiles from 73 cellular contexts across human, rhesus macaque and mouse. A transcript observed in multiple cellular contexts contributed one profile per context.

### Ablation experiments

Structural ablations retained TRACE's nucleotide embedding, prediction head, losses, chromosome partitions and optimization. LN-model used 12 standard pre-LayerNorm Transformer layers (model dimension, 384; attention heads, 16; feed-forward dimension, 768; dropout, 0.1) with sequence and padding-mask inputs only. Conv-model used 12 residual one-dimensional convolutional blocks (pre-LayerNorm; kernel width, 7; hidden dimension, 384; GELU;  $1 \times 1$  projection; dropout, 0.1). Its encoder parameter count was within 0.3% of LN-model. Both models ignored cellular-context and species inputs.

The cellular-context ablation used the TRACE backbone. ‘Zero’ training strategy forced expression to zero during training and validation; ‘Real’ training strategy supplied the measured 16,840-gene vector without perturbation. For ‘Augment’ training strategy, 10% of training samples received an exact zero vector; among the remainder, 30% were scaled by a value sampled uniformly from 0 to 1 after adding Gaussian noise (s.d., 0.15). Real versus Zero tested measured context, TRACE-Zero versus LN-model tested adaptive normalization without measured expression, and LN-model versus Conv-model tested global self-attention against local convolution.

Each strategy was trained on matched sets of 5, 22 or 40 human cellular environments; the 40-environment set combined 22 tissues and 18 common cell lines. All nine combinations were evaluated on the same chromosome-held-out transcripts from 26 uncommon cell lines absent from training. Generalization was summarized by nucleotide-profile Spearman  $\rho$ , periodicity-related performance, CDS-mean signal Spearman  $\rho$  and CDS-mean absolute error, weighting cellular environments equally.

Supplementary Table 2 | Model-ablation performance on the five-context dataset

| Model         | Best train loss<br>(epoch) | Best validation loss<br>(epoch) | Profile Spearman $\rho$<br>(epoch) | CDS-mean scale<br>Spearman $\rho$ (epoch) | CDS-mean MAE<br>(epoch) |
|---------------|----------------------------|---------------------------------|------------------------------------|-------------------------------------------|-------------------------|
| TRACE Zero    | 0.1187 (2)                 | 0.1093 (29)                     | 0.437 (44)                         | 0.612 (29)                                | 0.0981 (26)             |
| TRACE Real    | 0.0925 (50)                | 0.1067 (29)                     | 0.446 (47)                         | 0.649 (41)                                | 0.0953 (44)             |
| TRACE Augment | 0.0956 (50)                | 0.1062 (29)                     | 0.445 (47)                         | 0.645 (44)                                | 0.0957 (29)             |
| LN-model      | 0.0923 (1)                 | 0.1109 (42)                     | 0.427 (10)                         | 0.600 (44)                                | 0.1009 (11)             |
| Conv-model    | 0.0895 (1)                 | 0.1190 (5)                      | 0.399 (5)                          | 0.587 (40)                                | 0.1016 (38)             |

Table note. Values are the best values recorded across 50 epochs, with the corresponding epoch in parentheses. Training loss was computed on the training set; all other metrics used the human\_5c\_6k\_depth0.1\_cov0.1\_rpm1 validation set. Lower values are better for losses and CDS-mean MAE, whereas higher values are better for Spearman  $\rho$ . The best value in each column is shaded green. Because each metric was selected independently, values within a row can correspond to different checkpoints. Model names were standardized to the manuscript terminology: TRACE Augment corresponds to TRACE Mask+Interp. in the source file, LN-model to LN Transformer and Conv-model to Conv model (5c).

Supplementary Table 3 | Capability comparison of TRACE with prior translation models and ORF-identification tools

| Model           | Full-length RNA sequence | Variable-length input | CDS-annotation-free inference | Raw Ribo-seq-free inference | Single-nt ribosome-density output | Cell-context conditioning | Unified multi-species modeling | Direct TE prediction | Direct ORF identification |
|-----------------|--------------------------|-----------------------|-------------------------------|-----------------------------|-----------------------------------|---------------------------|--------------------------------|----------------------|---------------------------|
| TRACE           | Yes                      | Yes                   | Yes                           | Yes                         | Yes                               | Yes                       | Yes                            | Yes                  | Yes                       |
| Translatomer    | No                       | No                    | Yes                           | Yes                         | No                                | Yes                       | No                             | No                   | No                        |
| RiboMIMO        | No                       | Yes                   | No                            | Yes                         | No                                | No                        | No                             | No                   | No                        |
| Riboformer      | No                       | No                    | No                            | No                          | No                                | No                        | No                             | No                   | No                        |
| Optimus 5-Prime | No                       | No                    | No                            | Yes                         | No                                | No                        | No                             | Yes                  | No                        |
| RiboDecode      | No                       | No                    | No                            | Yes                         | No                                | Yes                       | No                             | Yes                  | No                        |
| RiboNN          | Yes                      | Yes                   | No                            | Yes                         | No                                | Yes                       | No                             | Yes                  | No                        |
| RiboTIE         | Yes                      | Yes                   | Yes                           | No                          | No                                | No                        | No                             | No                   | Yes                       |
| Ribo-TISH       | Yes                      | Yes                   | No                            | No                          | No                                | No                        | No                             | No                   | Yes                       |
| RibORF          | Yes                      | Yes                   | No                            | No                          | No                                | No                        | No                             | No                   | Yes                       |
| TranslationAI   | Yes                      | Yes                   | Yes                           | Yes                         | No                                | No                        | No                             | No                   | Yes                       |

Table note. Yes denotes direct support under the operational definitions below; No denotes that the capability is absent or is available only indirectly, through a fixed local window, a separate species-specific model, task-specific retraining, post hoc analysis or a required Ribo-seq input. All columns are phrased as positive capabilities, so Yes under raw Ribo-seq-free inference means that raw or processed Ribo-seq data are not required at inference. A No designation therefore indicates a difference in scope rather than lower overall model quality.

Definitions: Full-length RNA input includes either a complete transcript sequence (5' UTR, CDS and 3' UTR) or full-length transcript information reconstructed from a genome FASTA file and GTF/GFF annotation. Variable-length input excludes models restricted to a fixed local sequence window; padding or chunking of biologically variable transcripts is allowed. CDS-annotation-free inference requires no CDS, start-codon, reading-frame annotation or predefined 5'-UTR/CDS boundary at inference; Optimus 5-Prime is therefore classified as No. Raw Ribo-seq-free inference means that neither raw Ribo-seq reads nor processed ribosome-profiling signals, such as mapped RPF counts, P-site profiles or a matched control-condition ribosome-density track, are required as model inputs. Single-nt ribosome-density output refers specifically to a predicted ribosome-density value at each nucleotide, not TIS/TTS probabilities or codon-level profiles. Explicit cell-context conditioning requires a cell-environment representation supplied at inference; fixed cell-type output heads or matched control Ribo-seq tracks do not qualify. Unified multi-species modeling requires one jointly trained framework or checkpoint that models multiple species without species-specific retraining. Direct TE prediction requires TE or a directly corresponding translation-output quantity as a model output. Direct ORF identification requires explicit TIS/ORF discovery or translated-ORF classification.

Feature assignments were checked against the primary publications and official implementations summarized in TRACE\_prior\_work\_comparison.xlsx.